


Linear Regression



Outline

- Linear Regression with One Variable
- Multiple Linear Regression
- Polynomial Regression
- Regression Practical Considerations

Some of slides are based on the slides from

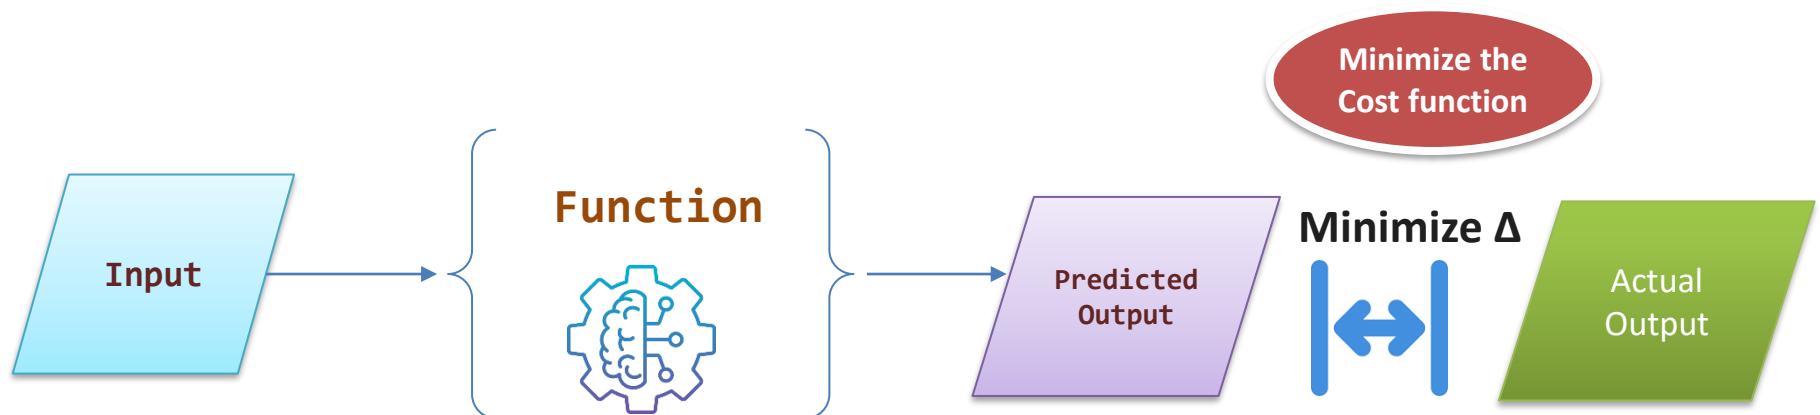
<https://www.deeplearning.ai/courses/machine-learning-specialization/>

Linear Regression with One Variable




ML: learn a **Function** that minimizes the cost

- Start with random function parameters
- Repeat intelligent guessing/approximation of the Function parameters such that the difference between the Predicted Output the Actual Output is reduced
 - i.e., minimize a Cost function a.k.a loss, or error function



Linear Regression with One Variable

Housing price prediction



Linear regression with one variable aka **Univariate linear regression**

- **Regression:** Predict continuous output value (e.g., price)
- Linear regression is used to predict the value of a variable based on the value of another variable(s)
- Linear regression **fits** a straight line that minimizes the discrepancies between predicted and actual output values

Training set of housing prices

Size in feet ² (x)	Price (\$) in 1000's (y)
2104	460
1416	232
1534	315
852	178
...	...

Notation:

m = Number of training examples

x = “input” variable / features

y = “output” variable / “target” variable

$(x^{(i)}, y^{(i)})$ – the i^{th} training example

$$x^{(1)} = 2104$$


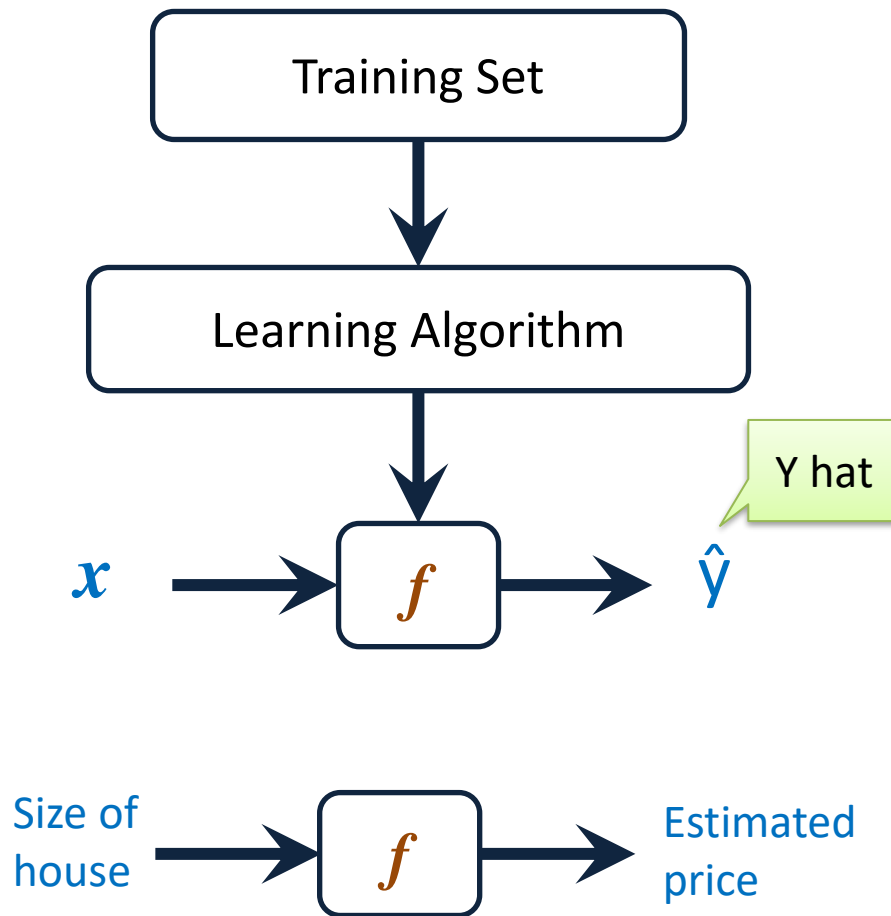
$$x^{(2)} = 1416$$

$$y^{(1)} = 460$$

How do we represent f ?

$$f(x) = wx + b$$

w, b are parameters (coefficients)
to learn from the training set





- Given a training set, **learn a function f** so that $f(x)$ is a “good” predictor for the corresponding value of y
- Find w, b parameters that **minimize** the error between predicted and actual values (i.e., minimize $\frac{1}{m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$ for all dataset instances)

Univariate Linear Regression - Model Representation


$$f(x) = wx + b$$

- w is the slope of the line
- b is the y-intercept of the line


How to choose w and b ?




$$w = 0$$
$$b = 1.5$$



$$w = 0.5$$
$$b = 0$$



$$w = 0.5$$
$$b = 1$$



Idea: Find w and b so that $f(x)$ is close to y for our training examples (x, y)

Find w, b :

$\hat{y}^{(i)}$ is close to $y^{(i)}$ for all $(x^{(i)}, y^{(i)})$

Cost (mean squared error)

Function:


$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2$$

Goal: minimize $J(w, b)$
 w, b

With m = number of training examples


Visualizing the Cost function

$$f(x) = wx$$



$$\begin{aligned} J(w) &= \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2 \\ &= \frac{1}{2m} (0^2 + 0^2 + 0^2) = 0 \end{aligned}$$


For simplicity, let us assume our optimization objective is to minimize $J(w)$, thus, $b = 0$



[slides\regression.xlsx](#)


Visualizing the Cost function

$$f(x) = wx$$




$$\begin{aligned} J(w) &= \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2 \\ &= \frac{1}{2m} ((0.5 - 1)^2 + (1 - 2)^2 + (1.5 - 3)^2) \\ &= \frac{1}{2 \times 3} (3.5) = \frac{3.5}{6} = 0.58 \end{aligned}$$

For simplicity, let us assume our optimization objective is to minimize $J(w)$, thus, $b = 0$



Visualizing the Cost function


$$f(x) = wx$$



$$J(w) = \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2$$



$$= \frac{1}{2m} (1^2 + 2^2 + 3^2)$$

$$= \frac{1}{2 \times 3} (14) = \frac{14}{6} = 2.3$$




$J(w)$ cost function
has **red bowl shape**
Convex function

For simplicity, let us assume our optimization objective is to minimize $J(w)$, thus, $b=0$



Visualizing the Cost function


$$f(x) = wx + b$$




$$f(x) = 0.06x + 50$$

- The fact that the cost function squares the error ensures that the 'error surface' is **convex** like a soup bowl.
- It will always have a minimum that can be reached by following the **gradient** (i.e., the slope)
- Minimizing the cost function yields optimal values of w and b


$$J(w, b)$$



$J(w, b)$ cost function
has **bowl shape**
Convex function



`08.regression\02_cost_function.py`




MSE has only 1 global minimum

Bowl shape
Convex function

Gradient Descent can be used to find the optimal w and b that minimizes that cost function



unsplash.com/photos/3m6vbzY69s4



Other cost functions may have multiple local minima

Gradient descent algorithm

Want to find w and b that minimize the cost function J $\underset{w,b}{\text{minimize}} J(w, b)$

1. Initialize the values of w and b to some arbitrary values (say 0, 0)
2. Calculate the predicted values of y using the current values of w and b
3. Calculate the gradients of the cost function with respect to w and b
4. Update the values of w and b using the gradients and a learning rate

The diagram shows the update equations for the weights w and bias b in the gradient descent algorithm. The first equation is $w = w - \alpha \frac{d}{dw} J(w, b)$, where α is circled in red and labeled "Learning Rate" with a green callout box. The term $\frac{d}{dw} J(w, b)$ is enclosed in a pink box and labeled "Derivative of the Cost Function w.r.t w " with a green callout box. The second equation is $b = b - \alpha \frac{d}{db} J(w, b)$.

$$w = w - \alpha \frac{d}{dw} J(w, b)$$
$$b = b - \alpha \frac{d}{db} J(w, b)$$

5. Repeat steps 2-4 until convergence (i.e., until the cost function converges to a minimum)

Gradient descent algorithm

Gradient descent utilizes the partial derivative of the cost function with respect to \mathbf{w} and \mathbf{b} to update \mathbf{w} and \mathbf{b} parameters

Repeat until convergence {

$$\mathbf{w} = \mathbf{w} - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) \cdot \mathbf{x}^{(i)}}_{\frac{\partial J}{\partial \mathbf{w}}}$$

$$\mathbf{b} = \mathbf{b} - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})}_{\frac{\partial J}{\partial \mathbf{b}}}$$


}

(simultaneously update \mathbf{w} and \mathbf{b})




Derivative 101

- Source <https://www.mathsisfun.com/calculus/derivatives-introduction.html>
- **Derivatives:** it is all about slope!



We can find an **average** slope between two points



Derivative = slope at a point

- Fill in this slope formula: $\frac{\Delta y}{\Delta x} = \frac{f(x+\Delta x) - f(x)}{\Delta x}$
- Simplify it as best we can
- Then make Δx shrink towards zero.

Example

The slope formula is: $\frac{f(x+\Delta x) - f(x)}{\Delta x}$

Use $f(x) = x^2$: $\frac{(x+\Delta x)^2 - x^2}{\Delta x}$


Expand $(x+\Delta x)^2$ to $x^2 + 2x \Delta x + (\Delta x)^2$: $\frac{x^2 + 2x \Delta x + (\Delta x)^2 - x^2}{\Delta x}$

Simplify (x^2 and $-x^2$ cancel): $\frac{2x \Delta x + (\Delta x)^2}{\Delta x}$

Simplify more (divide through by Δx): $2x + \Delta x$

Then, **as Δx heads towards 0** we get: $2x$

Result: the derivative of x^2 is $2x$




$$\frac{d}{dx} x^2 = 2x$$

In other words, the slope at x is $2x$

Interpretation of Derivative


- So what does $\frac{d}{dx}x^2 = 2x$ mean?
- It means that, for the function x^2 , the slope or "rate of change" at any point is $2x$
- So, when $x=2$ the slope is $2x = 4$
- Or when $x=5$ the slope is $2x = 10$, and so on



08.regression\
03_gradient_decent_x_square_animation.py

The Impact of Partial Derviative

- For simplicity, let us assume our optimization objective is to minimize $J(\mathbf{w})$, thus, $b = 0$



Learning Rate

$$w = w - \alpha \frac{dJ(w)}{dw}$$


$$= w - \alpha (\text{Positive Number})$$

Decrease w by a certain value

repeat until convergence {


$$w = w - \alpha \frac{dJ(w)}{dw}$$

}



The Impact of Partial Derviative

- For simplicity, let us assume our optimization objective is to minimize $J(\mathbf{w})$, thus, $b = 0$



$$w = w - \alpha \frac{dJ(w)}{dw}$$


$= w - \alpha (\text{Negative Number})$

Increase w by a certain value



The Impact of Partial Derviative

- For simplicity, let us assume our optimization objective is to minimize $J(\mathbf{w})$, thus, $b = 0$
 w, b




$$w = w - \alpha \frac{dJ(w)}{dw}$$

$$= w - \alpha (\text{Zero})$$

w remains the same, hence,
gradient descent *converges*

The Impact of Learning Rate




$$w = w - \alpha \frac{dJ(w)}{dw}$$

$$= w - (\text{Too Small Number}) \frac{dJ(w)}{dw}$$

What happens if α is too small?

w changes only a tiny bit on each step, hence, gradient descent *will render slow (will take more time to converge)*



$$w = w - \alpha \frac{dJ(w)}{dw}$$


$$= w - (\text{Too Large Number}) \frac{dJ(w)}{dw}$$

What happens if α is too large?

w changes a lot (and probably faster) on each step, hence, gradient descent *will potentially **overshoot the minimum** and, accordingly, fail to converge (or even diverge)*

The Impact of Learning Rate

We can set α between 0 and 1 (say, 0.1, or a little more or less, hence, not very small or very large)





$$w = w - \alpha \frac{dJ(w)}{dw}$$

α remains fix. As we approach the minimum, gradient descent will automatically start taking smaller steps (i.e., w will start changing at a slower pace because the derivative will become less steep)


Visualizing gradient descent algorithm

$$f(x) = \mathbf{w}x + b$$

$$J(\mathbf{w}, b)$$




The optimal values of \mathbf{w} and b are in the center of the inner most 'circle' of the Contour Plot




Visualizing gradient descent algorithm

$$f(x) = wx + b$$



$$J(w, b)$$




- The contour plot shows $J(w, b)$ over a range of w and b . The cost levels are represented by the rings
- The red crosses is the **path of gradient descent**. The path makes steady progress toward its goal. The initial steps are much larger than the steps near the goal.
- The optimal values of w and b are in the center of the inner most 'circle' of the Contour Plot

Different modes of gradient descent

- **Batch Gradient Descent (BGD):** Each iteration of gradient descent uses all the training examples
 - It provides a precise estimate of the gradient but can be computationally expensive, especially for large datasets
- **Stochastic Gradient Descent (SGD):** only a random sample from the dataset is used to compute the gradient in each iteration
 - It is computationally more efficient
 - However, it introduces more noise in the parameter updates, leading to more oscillations in the convergence path
- **Mini-Batch Gradient Descent:** divides the dataset into small mini-batches and computes the gradient using a mini-batch in each iteration
 - This approach combines the efficiency of stochastic gradient descent with the stability of batch gradient descent

Multiple Linear Regression



Linear Regression with multiple variables

Multiple features (variables)

Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

m

$n=4$

w_1
 w_2
 w_3
 w_4

Notation:

n = number of features

x_j = j^{th} feature

$\vec{x}^{(i)}$ = features of i^{th} training example

$x_j^{(i)}$ = value of feature j in i^{th} training example

Multiple Linear Regression - Model Representation

Univariate Linear Regression: $f(x) = wx + b$

Multiple Linear Regression:

$$f(x) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

Example $f(x) = 0.1x_1 + 4x_2 + 10x_3 + -2x_4 + 80$

↑ ↑ ↑ ↑ ↑
size #bedrooms #floors years base price

$$f_{\vec{w},b}(\vec{x}) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

$$\vec{w} = [w_1 \ w_2 \ w_3 \ \dots \ w_n]$$

parameters of the model

b is a number

vector $\vec{x} = [x_1 \ x_2 \ x_3 \ \dots \ x_n]$

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b = w_1x_1 + w_2x_2 + w_3x_3 + \dots + w_nx_n + b$$

↑
dot product

Multiple Linear Regression - Model Representation

Parameters and features

$$\vec{w} = [w_1 \ w_2 \ w_3] \quad n=3$$

b is a number

$$\vec{x} = [x_1 \ x_2 \ x_3]$$

linear algebra: count from 1

NumPy 

```
w = np.array([1.0, 2.5, -3.3])
```

```
b = 4
```

```
x = np.array([10, 20, 30])
```

code: count from 0

Without vectorization

$$f_{\vec{w},b}(\vec{x}) = w_1x_1 + w_2x_2 + w_3x_3 + b$$

```
f = w[0] * x[0] +  
     w[1] * x[1] +  
     w[2] * x[2] + b
```



Without vectorization

$$f_{\vec{w},b}(\vec{x}) = \left(\sum_{j=1}^n w_j x_j \right) + b$$

```
f = 0
```

```
for j in range(0, n):
```

```
    f = f + w[j] * x[j]
```

```
f = f + b
```



Vectorization

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$$

```
f = np.dot(w, x) + b
```



Multiple Linear Regression - Model Representation

Previous notation

Vector notation

Parameters

$$w_1, \dots, w_n \\ b$$

$$\vec{w} = [w_1 \quad \dots \quad w_n] \\ b$$

Model $f_{\vec{w},b}(\vec{x}) = w_1 x_1 + \dots + w_n x_n + b$

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$$

↑
dot product

Cost function $J(w_1, \dots, w_n, b)$

$$J(\vec{w}, b)$$

Gradient descent

repeat {

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(w_1, \dots, w_n, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w_1, \dots, w_n, b)$$

}

repeat {

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(\vec{w}, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(\vec{w}, b)$$

}

Vector representation of X, y and w

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$$

$$\vec{w} \cdot \vec{x} = \begin{bmatrix} w_0 \\ w_1 \\ \dots \\ w_n \end{bmatrix} \cdot \begin{bmatrix} x_0^{(0)} & x_1^{(0)} & \dots & x_n^{(0)} \\ x_0^{(1)} & x_1^{(1)} & \dots & x_n^{(1)} \\ \dots & \dots & \dots & \dots \\ x_0^{(m-1)} & x_1^{(m-1)} & \dots & x_n^{(m-1)} \end{bmatrix} = \begin{bmatrix} w_0 x_0^{(0)} + w_1 x_1^{(0)} + \dots + w_n x_n^{(0)} \\ w_0 x_0^{(1)} + w_1 x_1^{(1)} + \dots + w_n x_n^{(1)} \\ \dots \\ w_0 x_0^{(m-1)} + w_1 x_1^{(m-1)} + \dots + w_n x_n^{(m-1)} \end{bmatrix}$$

Dot product operation

$$\begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \cdot \begin{bmatrix} 2 \\ 4 \\ 6 \end{bmatrix} = \begin{bmatrix} \boxed{} & \boxed{} \\ \boxed{} & \boxed{} \\ \boxed{} & \boxed{} \end{bmatrix} = \begin{bmatrix} \boxed{} \\ \boxed{} \\ \boxed{} \end{bmatrix}$$

Dot product operation

$$\mathbf{a} \cdot \mathbf{b} = \mathbf{a}^T \mathbf{b} = \sum_{i=0}^{n-1} a_i b_i$$

$$\mathbf{a} \cdot \mathbf{b} = \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ a_3 \end{bmatrix} \begin{bmatrix} b_0 \\ b_1 \\ b_2 \\ b_3 \end{bmatrix} = [a_0 b_0 + a_1 b_1 + a_2 b_2 + a_3 b_4]$$

$$\mathbf{a}^T \mathbf{b} = \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ a_3 \end{bmatrix} \begin{matrix} a_1 & a_2 & a_3 \end{matrix} \begin{bmatrix} b_0 \\ b_1 \\ b_2 \\ b_3 \end{bmatrix} = [a_0 b_0 + a_1 b_1 + a_2 b_2 + a_3 b_4]$$

transform

Gradient Descent

One feature

repeat {

$$w = w - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)}) x^{(i)}$$

$$\frac{\partial}{\partial w} J(w, b)$$

$$b = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})$$

simultaneously update w, b

}

n features ($n \geq 2$)

repeat {

$$w_1 = w_1 - \alpha \frac{1}{m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)}) x_1^{(i)}$$

\vdots

$j=n$

$$w_n = w_n - \alpha \frac{1}{m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)}) x_n^{(i)}$$


$$b = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)})$$

simultaneously update

w_j (for $j = 1, \dots, n$) and b

}


Polynomial Regression



Polynomial Regression

- When the relationship between the set of features and the target variable is not linear then we need to use polynomial regression of some degree
 - The degree of the polynomial is usually a hyperparameter of the model

Polynomial Regression



$$\begin{aligned} f(x) &= w_1x_1 + w_2x_2 + w_3x_3 + b \\ &= w_1(\text{size}) + w_2(\text{size})^2 + w_3(\text{size})^3 + b \end{aligned}$$


$$x_1 = (\text{size})$$

$$x_2 = (\text{size})^2$$

$$x_3 = (\text{size})^3$$

- Polynomial regression is a special type of multiple regression whose independent variables are **powers** of variable X
- It is used to **approximate a curve** with unknown functional form
- For each additional power of X added to the model, the regression line will have one more bend

Polynomial Regression



$$f(x) = w_1(\text{size}) + w_2(\text{size})^2 + b$$

$$f(x) = w_1(\text{size}) + w_2\sqrt{(\text{size})} + b$$

Training Polynomial Regression Model

- Add **Polynomial Features** then train a Linear Regression Model
- For example, if we have two features then our new features will be: x_1 , x_2 , x_1x_2 , x_1^2 , x_2^2 , and the features matrix will be:

$$X = \begin{pmatrix} 1 & x_{11} & x_{12} & x_{11}x_{12} & x_{11}^2 & x_{12}^2 \\ 1 & x_{21} & x_{22} & x_{21}x_{22} & x_{21}^2 & x_{22}^2 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{n1} & x_{n2} & x_{n1}x_{n2} & x_{n1}^2 & x_{n2}^2 \end{pmatrix}$$


- Scikit-learn provides the transformer [PolynomialFeatures](#) that creates the features matrix consisting of all the polynomial combinations of the features up to a **specified degree**
- Then we can use the created features matrix to train a linear regression model

Polynomial Regression

- We can treat the powers of x : x, x^2, \dots, x^d , as distinct independent variables
 - Then, polynomial regression becomes a special case of multiple linear regression, since the model is still linear in the parameters that need to be estimated
- Therefore, we can find the optimal parameters \vec{w} using gradient descent






Regression Practical Considerations



Gradient descent in practice:

- Debugging Learning rate
- Feature Scaling
- Regularization
- Regression Evaluation

Learning rate: Gradient descent not working
Use smaller α



- For sufficiently small α , $J(w)$ should decrease on every iteration
- If α is too small, gradient descent can be slow to converge
- If α is too large: $J(w)$ may not decrease on every iteration; may not converge

To choose α , try


..., 0.001, 0.003, 0.01, 0.03, 0.1, 0.3, 1, ...

Feature Scaling: divide the input values by the range (i.e. the maximum value minus the minimum value) of the input variable, resulting in a new range of just 1.

The idea: Make sure features are on a similar scale. So that the gradient descent converges faster.


E.g. $x_1 = \text{size (0-2000 feet}^2\text{)}$

$x_2 = \text{number of bedrooms (1-5)}$



$$x_1 = \frac{\text{size (feet}^2\text{)}}{2000}$$

$$x_2 = \frac{\text{number of bedrooms}}{5}$$



Rule-of-thumb: Get every feature into approximately $-1 \leq x_i \leq 1$ range, $-0.5 \leq x_i \leq 0.5$, or other similar small ranges.

Mean normalization

- Replace x_i to make features have approximately zero

mean:

$$x_i = \frac{x_i - \mu_i}{s_i}$$

Where μ_i is the **average** of all the values for feature (i) (in the training set) and s_i is the standard deviation.

— e.g.,


$$x_1 = \frac{size - 1000}{400} \quad (\text{average size of the houses is 1000, and standard deviation is 400})$$

$$x_2 = \frac{\#bedrooms - 2}{3} \quad (\text{average \# of bedrooms is 2, and the standard deviation 3})$$

Regularization


- The problem of overfitting

Example: Linear regression (housing prices)




$$w_0 + w_1x$$

“underfit/high bias”



$$w_0 + w_1x + w_2x^2$$

“just right”



$$w_0 + w_1x + w_2x^2 + w_3x^3 + w_4x^4$$

“overfit/high variance”

Overfitting: If we have too many features, the learned function may fit the training set very well ($J(w) = \frac{1}{2m} \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)})^2 \approx 0$) but fail to generalize to new examples (predict prices on new examples)

Addressing overfitting:

x_1 = size of house

x_2 = no. of bedrooms

x_3 = no. of floors


x_4 = age of house

x_5 = average income in neighborhood

x_6 = kitchen size

⋮

x_{100}



Addressing overfitting:

Options:

1. Reduce number of features.

- Manually select which features to keep
- Use feature selection algorithm


2. Regularization.

- Keep all the features, but reduce magnitude/values of parameters w_j
- Works well when we have a lot of features, each of which contributes a bit to predicting y


Regularization

- Cost function

Intuition



$$w_0 + w_1x + w_2x^2$$



$$w_0 + w_1x + w_2x^2 + w_3x^3 + w_4x^4$$

Suppose we penalize and make w_3, w_4 really small

$$J(w) = \frac{1}{2m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2 + 1000 w_3^2 + 1000 w_4^2$$

$$w_3 \approx 0, \quad w_4 \approx 0$$

Regularization

Small values for parameters w_1, w_2, \dots, w_n

- “Simpler/smoother” function
- Less prone to overfitting


Housing:

- Features: x_1, x_2, \dots, x_{100}
- Parameters: w_1, w_2, \dots, w_{100}

$$J(w) = \frac{1}{2m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2$$

$$J(w) = \frac{1}{2m} \left[\sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n w_j^2 \right]$$

$$\min_w J(w)$$



Lasso (L1) Regularization

$$\frac{1}{2m} \sum_{i=1}^m (y - Xw)^2 + \alpha \sum_{j=1}^p |w_j|$$

Ridge (L2) Regularization

$$\sum_{i=1}^n (y - Xw)^2 + \alpha \sum_{j=1}^p w_j^2$$

Regularized linear regression

$$J(w) = \frac{1}{2m} \left[\sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n w_j^2 \right]$$

$$\min_w J(w)$$

Gradient descent

Repeat {

$$b = b - \alpha \frac{\partial}{\partial w_0} J(w)$$

$$w_j = w_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) x_j^{(i)} - \frac{\lambda}{m} w_j \right]$$

}

(j = 1, 2, 3, ..., n)

$\frac{\partial}{\partial w_j} J(w)$ "Regularized"

$$w_j = w_j \left(1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

$$1 - \alpha \frac{\lambda}{m} < 1$$

Regression Evaluation

- Performance measured by

- Mean Squared Error (MSE)

$$MSE = \frac{1}{n} \sum (y - \hat{y})^2$$

- Root-Mean-Squared-Error (RMSE)

$$RMSE = \sqrt{\frac{(y - \hat{y})^2}{n}}$$

- Mean-Absolute-Error (MAE)

$$MAE = \frac{1}{n} \sum |y - \hat{y}|$$

- ...others