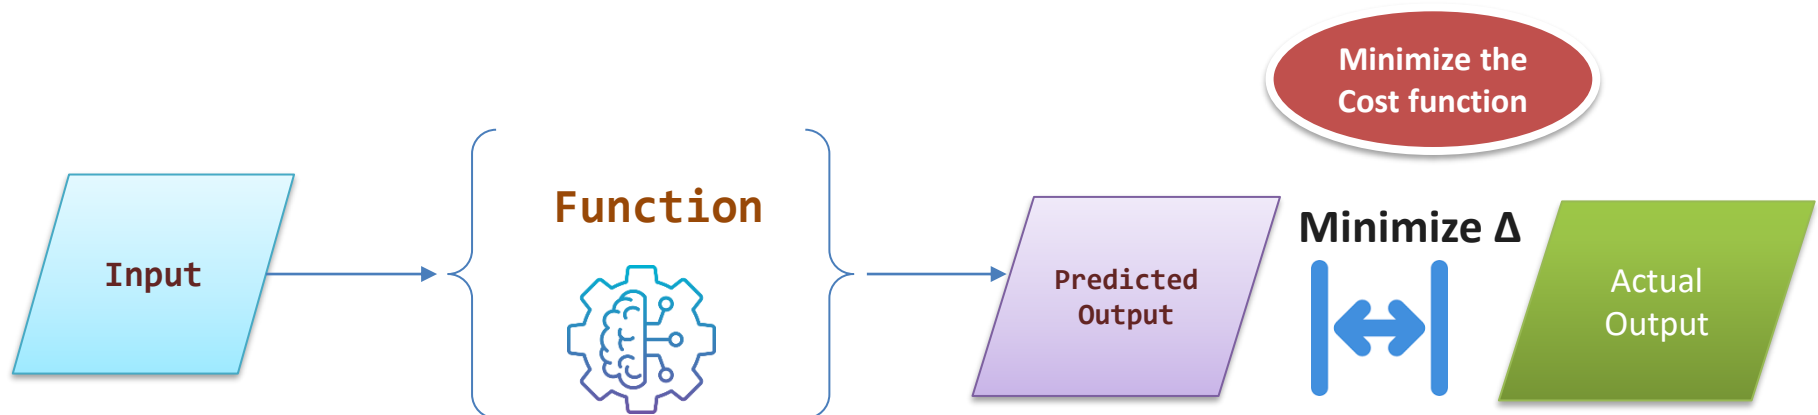



Regression

ML: learn a **Function** that minimizes the cost

- Start with random function parameters
- Repeat intelligent guessing/approximation of the Function parameters such that the difference between the Predicted Output the Actual Output is reduced
 - i.e., minimize a Cost function a.k.a loss, or error function



Linear Regression with One Variable



Regression: Predict continuous output value (price)

Training set of housing prices

Size in feet ² (x)	Price (\$) in 1000's (y)
2104	460
1416	232
1534	315
852	178
...	...

Notation:

m = Number of training examples

x = “input” variable / features

y = “output” variable / “target” variable

$$x^{(1)} = 2104$$

$$x^{(2)} = 1416$$


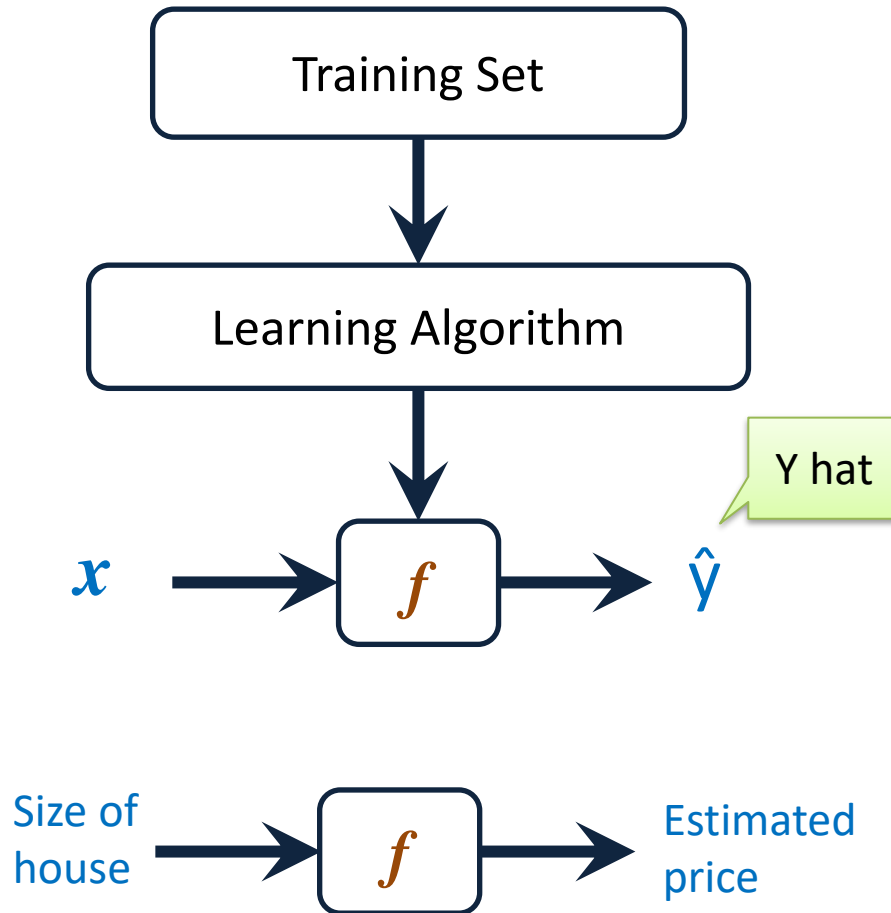
$$y^{(1)} = 460$$

$(x^{(i)}, y^{(i)})$ – the i^{th} training example

How do we represent f ?

$$f(x) = wx + b$$

w, b are parameters (coefficients)
to learn from the training set



Linear regression with one variable
Univariate linear regression



Given a training set, **learn a function f** so that $f(x)$ is a “good” predictor for the corresponding value of y (i.e. minimize the error between predicted and actual values)

Univariate Linear Regression - Model Representation


$$f(x) = wx + b$$

- w is the slope of the line
- b is the y-intercept of the line


How to choose w and b ?




$$w = 0$$
$$b = 1.5$$



$$w = 0.5$$
$$b = 0$$



$$w = 0.5$$
$$b = 1$$



Idea: Choose w and b so that $f(x)$ is close to y for our training examples (x, y)

Find w, b :

$\hat{y}^{(i)}$ is close to $y^{(i)}$ for all $(x^{(i)}, y^{(i)})$

Cost (mean squared error)

Function:

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2$$

Goal: minimize $J(w, b)$
 w, b

With m = number of training examples

Simplified

Function:

$$f(x) = wx + b$$


Parameters:

$$w, b$$

Cost Function:

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2$$

Goal: minimize $J(w, b)$
 w, b




$$f(x) = wx$$


$$w$$

$$J(w) = \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2$$

minimize $J(w)$
 w



$$f(x) = wx$$




$$J(w) = \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2$$


$$= \frac{1}{2m} (0^2 + 0^2 + 0^2) = 0$$

$$J(w)$$


(function of the parameter w)



$$f(x) = wx$$




$$J(w)$$



$$\begin{aligned} J(w) &= \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2 \\ &= \frac{1}{2m} ((0.5 - 1)^2 + (1 - 2)^2 + (1.5 - 3)^2) \\ &= \frac{1}{2 \times 3} (3.5) = \frac{3.5}{6} = 0.58 \end{aligned}$$

$$f(x) = wx$$





$$J(w) = \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2$$

$$= \frac{1}{2m} (1^2 + 2^2 + 3^2)$$


$$= \frac{1}{2 \times 3} (14) = \frac{14}{6} = 2.3$$

$$J(w)$$



$$f(x) = wx + b$$


(for fixed w, b this is a function of x)



$$f(x) = 50 + 0.06x$$


$$J(w, b)$$

(function of the parameters w, b)





$$f(x) = wx + b$$

(for fixed w, b this is a function of x)




$$J(w, b)$$

(function of the parameters w, b)




$$f(x) = wx + b$$

(for fixed w, b this is a function of x)




$$J(w, b)$$

(function of the parameters w, b)




$$f(x) = wx + b$$

(for fixed w, b this is a function of x)




$$J(w, b)$$

(function of the parameters w, b)




$$f(x) = wx + b$$

(for fixed w, b this is a function of x)




$$J(w, b)$$

(function of the parameters w, b)




$$f(x) = wx + b$$

(for fixed w, b this is a function of x)




$$J(w, b)$$

(function of the parameters w, b)




$$f(x) = wx + b$$

(for fixed w, b this is a function of x)




$$J(w, b)$$

(function of the parameters w, b)




$$f(x) = wx + b$$

(for fixed w, b this is a function of x)




$$J(w, b)$$

(function of the parameters w, b)




$$f(x) = wx + b$$

(for fixed w, b this is a function of x)




$$J(w, b)$$

(function of the parameters w, b)




$$f(x) = wx + b$$

(for fixed w, b this is a function of x)



$$J(w, b)$$

(function of the parameters w, b)



Linear Regression - Gradient decent


- Gradient decent


Have a cost function $J(\mathbf{w}, \mathbf{b})$

Want to find \mathbf{w} and \mathbf{b} $\min_{\mathbf{w}, \mathbf{b}} J(\mathbf{w}, \mathbf{b})$

Outline:

- Start with some \mathbf{w}, \mathbf{b} (say 0, 0)
- Keep changing \mathbf{w}, \mathbf{b} to reduce $J(\mathbf{w}, \mathbf{b})$
until we hopefully end up at a minimum





Gradient descent algorithm

1. Initialize the values of \mathbf{w} and \mathbf{b} to some arbitrary values
2. Calculate the predicted values of y using the current values of \mathbf{w} and \mathbf{b}
3. Calculate the gradients of the cost function with respect to \mathbf{w} and \mathbf{b}
4. Update the values of \mathbf{w} and \mathbf{b} using the gradients and a learning rate
5. Repeat steps 2-4 until convergence (i.e., until the cost function converges to a minimum)

Gradient descent algorithm

Repeat until convergence

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

Learning rate
Derivative

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

Simultaneously
update w and b

Correct: Simultaneous update

$$tmp_w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$tmp_b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

$$w = tmp_w$$

$$b = tmp_b$$

Incorrect

$$tmp_w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$


$$w = tmp_w$$

$$tmp_b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$


$$b = tmp_b$$

Derivative 101

- Source <https://www.mathsisfun.com/calculus/derivatives-introduction.html>
- **Derivatives:** it is all about slope!





We can find
an **average** slope between
two points



But how do we find the slope at a point?

- There is nothing to measure!
slope $0/0 = \text{????}$
 - But with derivatives we use a small difference ...
- ... then have it shrink towards zero

- Fill in this slope formula: $\frac{\Delta y}{\Delta x} = \frac{f(x+\Delta x) - f(x)}{\Delta x}$
- Simplify it as best we can
- Then make Δx shrink towards zero.



Derivative Example - $f(x) = x^2$

The slope formula is: $\frac{f(x+\Delta x) - f(x)}{\Delta x}$

Use $f(x) = x^2$: $\frac{(x+\Delta x)^2 - x^2}{\Delta x}$

Expand $(x+\Delta x)^2$ to $x^2 + 2x \Delta x + (\Delta x)^2$: $\frac{x^2 + 2x \Delta x + (\Delta x)^2 - x^2}{\Delta x}$

Simplify (x^2 and $-x^2$ cancel): $\frac{2x \Delta x + (\Delta x)^2}{\Delta x}$

Simplify more (divide through by Δx): $2x + \Delta x$

Then, **as Δx heads towards 0** we get: $2x$

Result: the derivative of x^2 is $2x$

$$\frac{d}{dx} x^2 = 2x$$


- In other words, the slope at x is $2x$

Interpretation of Derivative

- So what does mean?

$$\frac{d}{dx}x^2 = 2x$$

- It means that, for the function x^2 , the slope or "rate of change" at any point is **$2x$**
- So when **$x=2$** the slope is **$2x = 4$**
- Or when **$x=5$** the slope is **$2x = 10$** , and so on



Gradient descent algorithm



repeat until convergence {
 $w = w - \alpha \frac{d}{dw} J(w)$
}



unsplash.com/photos/3m6vbzY69s4


Simplified

$$w = w - \alpha \frac{d}{dw} J(w)$$




$$w = w - \alpha \frac{d}{dw} J(w)$$

If α is too small, gradient descent can be slow




If α is too large, gradient descent can overshoot the minimum. It may fail to converge, or even diverge



Gradient descent can converge to a local minimum, even with the learning rate α fixed

$$w = w - \alpha \frac{d}{dw} J(w)$$

As we approach a local minimum, gradient descent will automatically take smaller steps (the slope gets smaller). So, no need to decrease α over time



Linear Regression with One Variable

Linear regression model **Cost function**

$$f_{w,b}(x) = wx + b$$

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2$$

Gradient descent algorithm

repeat until convergence {

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b) \quad \rightarrow \quad \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})x^{(i)}$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b) \quad \rightarrow \quad \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})$$


}

“Batch” Gradient Descent

“Batch”: Each step of gradient descent uses all the training examples, m

	x size in feet ²	y price in \$1000's
(1)	2104	400
(2)	1416	232
(3)	1534	315
(4)	852	178
...
(47)	3210	870

$$\sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2$$



Linear Regression with multiple variables

Multiple features (variables)

Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

Notation:

n = number of features

x_j = j^{th} feature

$\vec{x}^{(i)}$ = features of i^{th} training example

$x_j^{(i)}$ = value of feature j in i^{th} training example

Hypothesis:

Previously: $f(x) = wx + b$

Now: Multivariate linear regression.

$$f(x) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

Example

$$f(x) = 0.1x_1 + 4x_2 + 10x_3 + -2x_4 + 80$$

↑ ↑ ↑ ↑ ↑
size #bedrooms #floors years base price

Parameters: w_0, w_1, \dots, w_n

Cost function:

$$J(w_0, w_1, \dots, w_n) = \frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2$$

Gradient descent:

repeat {

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(w_0, \dots, w_n)$$

} (simultaneously update for every $j = 0, \dots, n$)

Gradient Descent

Previously ($n=1$):

Repeat {

$$w = w - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})}_{\frac{\partial}{\partial w} J(w)}$$

$$b = b - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) \cdot x^{(i)}}_{\frac{\partial}{\partial b} J(w)}$$

(simultaneously update w_0 and w_1)

New algorithm ($n \geq 1$):

Repeat {

$$w_j = w_j - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) x_j^{(i)}}_{\frac{\partial}{\partial w_j} J(w)}$$

(simultaneously update w_j for
 $j = 0, \dots, n$)

}

$$w_0 = w_0 - \alpha \frac{1}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$w_1 = w_1 - \alpha \frac{1}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) x_1^{(i)}$$

$$w_2 = w_2 - \alpha \frac{1}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) x_2^{(i)}$$

....

Formula to update **b** remains the same

Gradient descent in practice:


- Feature Scaling
- Regularization
- Regression Evaluation

Feature Scaling: divide the input values by the range (i.e. the maximum value minus the minimum value) of the input variable, resulting in a new range of just 1.

The idea: Make sure features are on a similar scale. So that the gradient descent converges faster.


E.g. $x_1 = \text{size (0-2000 feet}^2\text{)}$

$x_2 = \text{number of bedrooms (1-5)}$



$$x_1 = \frac{\text{size (feet}^2\text{)}}{2000}$$

$$x_2 = \frac{\text{number of bedrooms}}{5}$$



Rule-of-thumb: Get every feature into approximately a $-1 \leq x_i \leq 1$ range, $-0.5 \leq x_i \leq 0.5$, or other similar small ranges.

Mean normalization

- Replace x_i to make features have approximately zero mean (Do not apply to $x_0 = 1$):

$$x_i := \frac{x_i - \mu_i}{s_i}$$

Where μ_i is the **average** of all the values for feature (i) (in the training set) and s_i is the range of values ($max - min$), or s_i is the standard deviation.

– e.g.,

$$x_1 = \frac{size - 1000}{2000} \quad (\text{average size of the houses is 1000, and ranges from 0 to 2000})$$


$$x_2 = \frac{\#bedrooms - 2}{4} \quad (\text{average \# of bedrooms is 2, and the range is from 1 to 5})$$

$$-0.5 \leq x_1 \leq 0.5, -0.5 \leq x_2 \leq 0.5,$$

Regularization


- The problem of overfitting

Example: Linear regression (housing prices)




$$w_0 + w_1x$$

“underfit/high bias”



$$w_0 + w_1x + w_2x^2$$

“just right”



$$w_0 + w_1x + w_2x^2 + w_3x^3 + w_4x^4$$

“overfit/high variance”

Overfitting: If we have too many features, the learned hypothesis may fit the training set very well ($J(w) = \frac{1}{2m} \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)})^2 \approx 0$) but fail to generalize to new examples (predict prices on new examples).

Addressing overfitting:

x_1 = size of house

x_2 = no. of bedrooms

x_3 = no. of floors


x_4 = age of house

x_5 = average income in neighborhood

x_6 = kitchen size

⋮

x_{100}



Addressing overfitting:


Options:

1. Reduce number of features.
 - Manually select which features to keep.
 - Use feature selection algorithm.
2. Regularization.
 - Keep all the features, but reduce magnitude/values of parameters w_j
 - Works well when we have a lot of features, each of which contributes a bit to predicting y .


Regularization

- Cost function

Intuition



$$w_0 + w_1x + w_2x^2$$



$$w_0 + w_1x + w_2x^2 + w_3x^3 + w_4x^4$$

Suppose we penalize and make w_3 , w_4 really small

$$J(w) = \frac{1}{2m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2 + 1000 w_3^2 + 1000 w_4^2$$

$$w_3 \approx 0, \quad w_4 \approx 0$$

Regularization

Small values for parameters w_0, w_1, \dots, w_n

- “Simpler/smooth” hypothesis
- Less prone to overfitting


Housing:

- Features: x_1, x_2, \dots, x_{100}
- Parameters: $w_0, w_1, w_2, \dots, w_{100}$

$$J(w) = \frac{1}{2m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2$$

$$J(w) = \frac{1}{2m} \left[\sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n w_j^2 \right]$$


$$\min_w J(w)$$



In regularized linear regression, we choose w to minimize

$$J(w) = \frac{1}{2m} \left[\sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n w_j^2 \right]$$

What if λ is set to an extremely large value (perhaps far too large for our problem, say $\lambda = 10^{10}$)?



$$w_1 \approx 0, \quad w_2 \approx 0, \quad w_3 \approx 0, \quad w_4 \approx 0$$

$$w_0 + w_1x + w_2x^2 + w_3x^3 + w_4x^4$$

Regularized linear regression

$$J(w) = \frac{1}{2m} \left[\sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n w_j^2 \right]$$

$$\min_w J(w)$$

Gradient descent

Repeat {

$$w_0 = w_0 - \alpha \frac{\frac{\partial}{\partial w_0} J(w)}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$w_j = w_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) x_j^{(i)} - \frac{\lambda}{m} w_j \right]$$

} $(j = \cancel{0}, 1, 2, 3, \dots, n)$ $\frac{\partial}{\partial w_j} J(w)$ "Regularized"

$$w_j := w_j \left(1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

$$1 - \alpha \frac{\lambda}{m} < 1$$

Regression Evaluation

- Performance measured by

- Mean Squared Error (MSE)

$$MSE = \frac{1}{n} \sum (y - \hat{y})^2$$

- Root-Mean-Squared-Error (RMSE)

$$RMSE = \sqrt{\frac{(y - \hat{y})^2}{n}}$$

- Mean-Absolute-Error (MAE)

$$MAE = \frac{1}{n} \sum |y - \hat{y}|$$

- ...others