


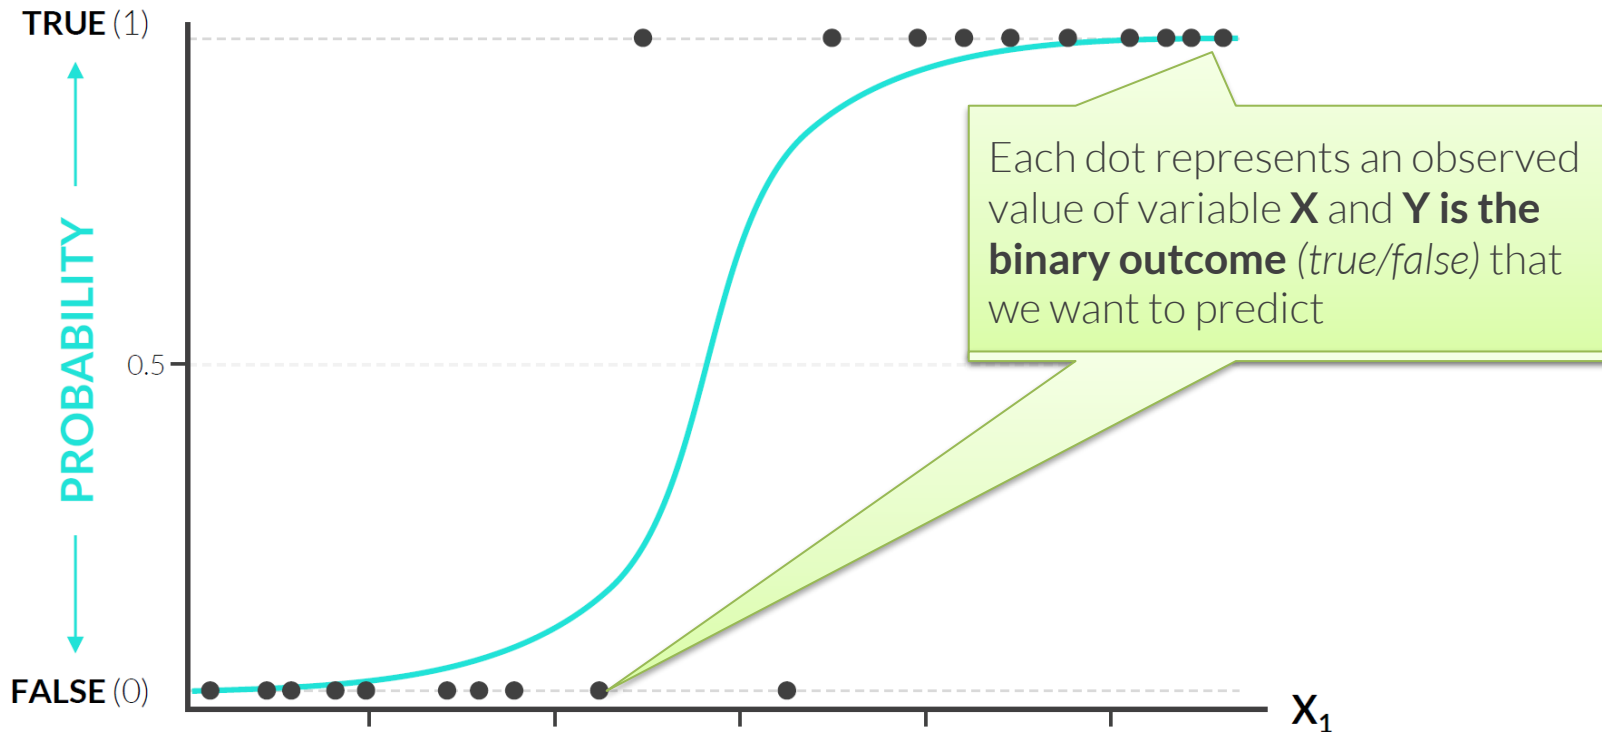
Logistic Regression



Logistic Regression


- Logistic Regression is a **classification** technique used to predict the probability of a binary (true/false) outcome
 - Although it has the word "regression" in its name, logistic regression is not used for predicting numeric variables
- Example use cases:
 - Classifying spam emails or fraudulent credit card transactions
 - Determining whether to serve a particular ad to a website visitor
- Logistic regression forms an S-shaped curve between 0 and 1 which represents the probability of a TRUE outcome for any given value of X
- The **cost function** is used to measure how accurately a model predicts outcomes, and is used to optimize the "shape" of the curve

Logistic Regression




- Logistic regression plots the **best-fitting curve between 0 and 1**, which tells us the probability of Y being TRUE for any given value of X_1

Logistic Regression - Example




- Using this model, we can classify unobserved values of X_1 (number of recipients) to predict the probability that Y is true or false (i.e., the probability that an email is spam)
- In practice a threshold of 0.5 is a common a decision point for logistic models ($P > 0.5$ means Y is predicted to be True)

Is 50% always the right decision point for logistic regression models?



- When the cost of a **false positive** (incorrectly predicting a TRUE outcome) is high we may increase the threshold (e.g., 90%), to avoid classifying legit emails as spam

(Incorrectly mark few spam emails as “not spam” is not a big deal)





- When the risk of a false negative (incorrectly predicting an employee will stay) is high we may decrease the threshold (e.g., 10%), to correctly predict more cases where someone is likely to quit

(It's easier to train and retain an employee than hire a new one, so the cost of a false negative -incorrectly predicting an employee will stay- is high)

sigmoid function

- The sigmoid function, also known as the logistic function, is a function that maps any real number into a range between 0 and 1
- The sigmoid function, denoted with the Greek letter sigma (σ), is defined as:

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



Logistic Regression - Model Representation

- Given a training set, learn a function \mathbf{f} so that $\mathbf{f}(\mathbf{x})$ is a “good” predictor for the corresponding value of y
 - Learn the weights (w) and bias (b) given inputs (x)

$$z = \left(\sum_{i=1}^n w_i x_i \right) + b \quad \xrightarrow{\text{Vector Notation}} \quad z = \mathbf{w} \cdot \mathbf{x} + b$$

W and X are vectors

- Furthermore, to get your prediction, you must apply the **sigmoid** function

$$\mathbf{f}(\mathbf{x}) = \hat{y} = \sigma(z) = \frac{1}{1+e^{-z}}$$

Learning a Logistic Regression Model

- Learn $\mathbf{w} = [w_1, \dots, w_m]$ and \mathbf{b} by minimizing the following cost function:

$$J = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

- $y^{(i)}$ is the actual label (0 or 1) for the i th training example.
- $\hat{y}^{(i)}$ is the predicted probability that the i th example belongs to class 1, given input $x^{(i)}$.
- The cost function (also known as a loss function) used is Binary cross entropy
 - It measures the difference between the predicted probabilities $\hat{\mathbf{y}}$ and the actual binary labels \mathbf{y}
 - This cost function essentially penalizes the model for predicting probabilities far from the actual labels
 - If the actual label is 1, it penalizes the model more for predicting a probability close to 0 (as given by $\log(1 - \hat{y}^{(i)})$ term), and vice versa

Gradient descent algorithm

Want to find w and b that minimize the cost function J minimize $J(w, b)$
 w, b

1. Initialize the values of w and b to some arbitrary values (say 0, 0)
2. Calculate the predicted values of y using the current values of w and b
3. Calculate the gradients of the cost function with respect to w and b
4. Update the values of w and b using the gradients and a learning rate

The diagram shows the update equations for the weights w and bias b in the gradient descent algorithm. The equation for w is $w = w - \alpha \frac{d}{dw} J(w, b)$, where α is circled in red and labeled "Learning Rate" with a callout box. The derivative term $\frac{d}{dw} J(w, b)$ is enclosed in a pink box and labeled "Derivative of the Cost Function w.r.t w " with a callout box. Below it, the equation for b is $b = b - \alpha \frac{d}{db} J(w, b)$.

$$w = w - \alpha \frac{d}{dw} J(w, b)$$
$$b = b - \alpha \frac{d}{db} J(w, b)$$

5. Repeat steps 2-4 until convergence (i.e., until the cost function converges to a minimum)

Gradient descent algorithm

Gradient descent utilizes the partial derivative of the cost function with respect to \mathbf{w} and \mathbf{b} to update \mathbf{w} and \mathbf{b} parameters

Repeat until convergence {

$$\mathbf{w} = \mathbf{w} - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) \cdot \mathbf{x}^{(i)}}_{\frac{\partial J}{\partial \mathbf{w}}}$$

Learning rate α , controls how big a step we take when we update \mathbf{w} and \mathbf{b}

$$\mathbf{b} = \mathbf{b} - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})}_{\frac{\partial J}{\partial \mathbf{b}}}$$

}

(simultaneously update \mathbf{w} and \mathbf{b})

